



UNIVERSITY OF AMSTERDAM

User manual for UvA AI Chat

Version 1.1

Published 2nd of October 2025



About this manual

This manual for UvA AI Chat is designed as a reference guide where you can look up information about the various features and settings of UvA AI Chat. If you're using UvA AI Chat for the first time -or if you don't have much experience with generative AI - we recommend starting with the "[Getting Started with the UvA AI Chat](#)" e-learning. This e-learning module walks you through all the basic features step by step and shows examples of how you can use the tool. Once you're familiar with UvA AI Chat, you can use this manual as a reference whenever you need it. You can also use the "UvA AI Chat Ondersteuning" persona within UvA AI Chat, which is based on this guide.

Because AI is evolving rapidly, this manual will also be regularly updated. The latest features and applications of UvA AI Chat will be included in future versions. You'll always find the most recent version of this manual in UvA AI Chat by clicking the question mark icon on the bottom left.

For questions about this guide and UvA AI Chat, please contact ai-tlc@uva.nl.

This manual is developed by TLC Central (see tlc.uva.nl)

Version 1.1

Published 2nd of October 2025



Part 1: Before you start	4
1.1 Responsible use	4
1.2 Your first step: logging in.....	4
Part 2: Core Features in practice	5
2.1 Communicating effectively: the art of prompting.....	5
2.2 “My prompts”: Your personal collection of instructions.....	7
2.3 Uploading and analysing documents.....	7
2.4 Generating and analyzing images	9
2.5 Artifacts.....	10
2.6 Study Mode.....	11
Part 3: Personalization and collaboration.....	12
3.1 Further configuration	12
3.2 Customized interaction: working with personas	12
3.3 “Groups”: collaborating and sharing with others	14
3.4 “My Projects”: your organized workspace	14
Part 4: Advanced usage	15
4.1 Choosing the right engine	15
4.2 Expanding functionality with extensions.....	16
Frequently Asked Questions (FAQ).....	17
Glossary.....	19

Part 1: Before you start

1.1 Responsible use

It's essential to understand that UvA AI Chat, like any other AI tool, is not infallible. The model can present mistakes, inaccuracies, biases, or incorrect information as if it were fact. Never rely blindly on the generated output. Always verify the results carefully and apply critical judgement each time you use it. Also keep in mind that using AI consumes both energy and water.

Practical example of responsible use:

Suppose you ask the AI to summarize a scientific article you've uploaded. The AI produces a concise summary.

- **Irresponsible use:** You copy the summary directly and use it in your own work without any verification.
- **Responsible use:** You read the summary generated by the AI and compare it carefully to the original article. You check whether the main arguments, methodology, and conclusions are represented correctly. You use the AI output as a tool to help you grasp the article faster, but you base your own work on the actual source text.

Part of using generative AI responsibly is being aware of the possible downsides. When using AI, you should be aware of the risks of cognitive offloading. Cognitive offloading refers to the process of transferring cognitive processes to an external tool, such as an AI. Especially in an educational setting it is important to be aware that continuous reliance on AI for a particular task might hinder you from developing knowledge and skills that would benefit you in the long run. Research shows that using AI frequently but unreflectively can lessen critical thinking abilities.

To use AI effectively and responsibly, we need to develop our *AI literacy*. Among other things, this means understanding how to use AI tools well, like asking the right questions, and carefully checking what the AI produces. By developing our AI literacy we can ensure that we make use of AI in such a way that it supports our work, learning, and research in a valuable manner, while mitigating the risks the use of AI poses.

1.2 Your first step: logging in

Open AI Chat in your webbrowser. Enter your UvA login details - this means your email address and password linked to your UvA account. For additional security, two-factor authentication (2FA) is required; you'll need to enter a verification code sent to your mobile device. After successful authentication, you'll have access to the UvA AI Chat dashboard and can use all its features.

Part 2: Core Features in practice

2.1 Communicating effectively: the art of prompting

The most important skill for making the best use of UvA AI Chat is writing effective prompts. The quality of the output you receive from the AI is directly tied to the quality of the instructions you provide.

Definition of a prompt

A prompt is an instruction or question you give to the AI to generate a specific response or output. It's the foundation of your interaction with the AI. The clarity and level of detail in your prompt have a direct impact on the quality of the answer. A well-crafted prompt acts as a clear assignment, while a vague prompt typically results in a generic or unusable response.

While crafting a strong prompt is important, there's no need to worry if your first attempt doesn't work perfectly. The real strength of the AI lies in its ability to carry on a conversation, so if the output isn't what you expected, you can always ask the AI to adjust its answer. It's also helpful to ask the AI to help you improve your prompt.

Structure of an effective prompt

Prompting effectively means making your prompts clear, specific, and relevant. To do this, state clearly what you want the AI to do, including any nuances or specific requirements. Use concrete language, avoid vagueness, and provide context where needed to help the AI understand your question and generate accurate answers. A well-structured prompt usually contains three main elements:

1. **Context (introduction):** Start with a brief introduction or background so the AI understands what the prompt is about. Who are you? What's the purpose of the task? Who is the output intended for?
2. **Question or instruction:** Clearly state your question or give a specific, unambiguous task. Use complete sentences and be explicit about the action you want the AI to take.
3. **Specific details or criteria:** Add relevant details, constraints, or criteria the AI should consider in its response. This could include the desired format (table, bullet points, continuous text), tone (formal, informal, creative), length (number of words), or specific points to address.

A key tip is to phrase instructions positively. Avoid exclusions (like "don't use difficult words"), because the AI tends to focus on any concept you mention. Instead, describe what you do want ("use simple, accessible language").

Practical example of positive formulation

- **Negative:** "Write an introduction for my essay, but don't use jargon and don't make it longer than 200 words."
- **Positive:** "Write an accessible introduction of around 200 words for my essay on Italo Calvino, aimed at a first-year student. Use clear, simple language and define the main concepts."

Table 1: Examples of positive formulation

The table below illustrates the difference between a vague prompt and a specific, well-structured prompt for various academic tasks.

Task	Ineffective Prompt	Effective Prompt
Writing an email	"Write an email to my professor"	“Context: I need to email my professor, Dr. Jansen, about an extension for my paper on the Eighty Years’ War. Instruction: Draft a formal and respectful email in which I request a two-day extension. Criteria: Mention that the reason is unforeseen personal circumstances (without going into detail) and suggest a new submission date.”
Summarizing text	“Summarize this document”	“Context: The uploaded document is a scientific article about the impact of AI on the labor market. Instruction: Generate a summary in bullet points. Criteria: Focus on the research methodology, main findings, and the authors’ conclusions. The summary should be in Dutch and no longer than 250 words.”
Brainstorming	“Give me ideas for an argument”	Context: I am a sociology researcher and want to write an essay within the theme of social inequality. Instruction: Generate five concrete standpoints I could write an essay about. Criteria: Each standpoint should be original, empirically grounded, and relevant to contemporary Dutch society.
Requesting feedback	“Check my text”	Context: The attached file is the draft introduction of my research proposal. The audience is researchers in my field. Instruction: Provide constructive feedback on the text. Criteria: Analyze the text on three points: 1) clarity of the problem statement, 2) logical structure of the argument, and 3) academic writing style. Give concrete examples from the text and suggest specific improvements.

2.2 “My prompts”: Your personal collection of instructions

The “My prompts” is a feature that helps you work more efficiently by reusing effective prompts. You’ll find the “My prompts” section via the book icon in the left sidebar of UvA AI Chat.

Using standard prompts

“My prompts” includes a collection of predefined prompts for common tasks. Examples include “feedback on your writing” and the “multiple choice question generation.” These standard prompts are designed by experts and already have a strong structure. You can select a standard prompt and easily adapt it to your specific needs, allowing you to get fast, effective results.

Saving and reusing your own prompts

If you notice you’re repeating certain tasks or instructions, you can save your own prompts in “My prompts”. This is especially useful for complex, recurring assignments.

Practical example: Saving a custom prompt

Suppose you frequently write academic texts in English and want them checked for formal writing style. You can draft a highly effective prompt for this and save it for future use.

1. Write your effective prompt.
2. Go to “My prompts” and select “Add prompt”.
3. Give your prompt a clear name, like “Academic English Check”.
4. Paste your prompt in the text field: “Analyze the attached English text. Act as an experienced editor for an academic journal. Identify and correct sentences that are too informal for scientific publication. Replace colloquial language with formal alternatives, check for consistency in terminology, and suggest ways to vary sentence structure for better readability.”
5. Save the prompt. It’s now available for easy reuse in future conversations.

Want to quickly reuse a previously saved prompt? Just click the book icon at the bottom of the text box in UvA AI Chat to see a list of your recent prompts and insert the one you need instantly.

Letting the AI search the internet

UvA AI Chat has a 'knowledge cutoff'; its knowledge is not up-to-date. To have the AI work with recent, online information, you can enable the internet search function. To do this, click on the 'swiss army knife' icon in the top right of a chat window and check 'Internet Search'. After you have activated the function, it is important that you state in your prompt that the AI should use the internet. For example, start your question with: "Search the internet for...". Please keep in mind that this functionality is still in development and does not always work flawlessly.

2.3 Uploading and analyzing documents

One of the most powerful features of UvA AI Chat is the ability to upload your own documents and use them as a contextual basis for your queries. This allows the AI to work with the specific information you provide, rather than just its general knowledge base. This makes the chatbot more powerful and reliable as it can give more accurate and well-supported answers based on your chosen sources.

Uploading documents

With UvA AI Chat, uploading documents is simple. Go to the text entry field at the bottom of the screen and click the paperclip icon. Then select the files you want to add from your computer. Once uploaded, the documents are immediately available for further use within the current conversation.

Supported file formats

UvA AI Chat supports a range of common file formats, including PDF, DOCX (Word), XLS (Excel), PPT (PowerPoint), and CSV. This makes it possible to upload and work with different document types, ranging from research articles to datasets, directly in your interactions with the AI.

Note: Large, unusual, or password-protected files may not always process correctly. If a document is very large (e.g., 100+ pages), the AI is essentially searching for a needle in a haystack.

Processing options and workflow

Once uploaded, the AI won't do anything with the document until you give an instruction. The document acts as a knowledge source for the AI. You can have texts analyzed, summarized, translated, or have the AI retrieve specific information from the document. The key is combining document upload with a precise prompt.

Practical example of integration in your workflow:

A researcher is working on a literature review and has a 50-page PDF report that's relevant to her research.

1. **Uploading:** She clicks the paperclip and uploads the PDF report to the chat.
2. **Prompting:** Next, she asks a specific question about the document's content. An effective prompt could be: "Analyze the recently uploaded report and identify all sections related to sustainable policy. Summarize the key recommendations in bullet points, and for each recommendation, cite the page number where it can be found."

With this approach, the researcher saves time and receives a focused, actionable analysis of the uploaded text, ready for use in her own work.

Retrieval-Augmented Generation (RAG)

The technology by which uploading and analyzing documents functions is called Retrieval-Augmented Generation (RAG). This works by augmenting the language model's generation with information retrieved from your uploaded documents. When you submit a prompt, the system first compares your query to your documents and then retrieves pieces of information from your documents that seem relevant for your prompt. It can take a maximum of 10 pieces of information. These snippets are then added to your prompt as context and processed by the language model (LLM) to generate a more accurate and well-supported answer. This is immensely powerful, as it allows the AI to base its answer on specific contextual information you choose. However, it also has limitations; for example, it is not designed to summarize very large texts (it can only take 10 snippets).

2.4 Generating and analysing images

In addition to text-based tasks, UvA AI Chat can also generate images based on your descriptions. This is useful for creating illustrations for presentations, visualizing concepts, or for creative projects. To generate an image, simply ask for one in the chat's text box. As with textual prompts, the more detailed and specific your description, the better the result will fit your expectations. As always, you can ask AI to help you refine your prompt. Note that the image generator in UvA AI Chat is not yet reliable for creating images containing text.

Practical example of a detailed image prompt:

Instead of asking something like “make a picture of a flower,” you’ll get far better results with a detailed prompt.

- **Too general:** “Generate a tulip”.
- **Detailed:** “Generate an image in the style of a 19th-century botanical illustration. The image should show a tulip with clear, fine lines highlighting the details of the petals, stem, and leaves. The background should resemble parchment with a light texture”.

Analyzing images

Use the photo icon at the bottom right of the text box to add images. After uploading, the AI can describe and interpret the content. Always include a prompt explaining what you want the AI to analyze. Note: the AI cannot always accurately interpret text, charts, or tables contained within images.

2.5 Artifacts

The "Artifacts" feature in UvA AI Chat is a powerful tool designed to support collaborative work with the AI. An artifact is a separate window (with text, code, or visualizations) that both you and the AI can edit. This allows for easy collaboration as you can ask the AI to edit the artifact, and you can select pieces of text in the artifact and tell the AI to specifically edit those. Additionally, you can edit the text manually. This makes the AI more into a creative partner rather than just a text generator.

How to use artifacts:

Activate the feature: To begin, you must select the "Artifact Creation" feature within the UvA AI Chat interface (click the slider menu). This signals to the AI that you are looking for a output in a separate window for further editing.

Provide a clear instruction: In the chat box, enter an instruction that clearly defines the topic and the type of artifact you need. For example, you could prompt the AI with: "create an artifact comparing Keynesian and Neoclassical economic theories."

Review the generated output: The AI will then produce an artifact based on your request. In the example of economic theories, the output could be an overview of the key principles, assumptions and political implications.

Collaborate on revisions: The platform allows you to refine the artifact collaboratively. You can make direct edits in the dedicated text or code window, or you can give the AI further instructions in the chat to modify or expand a specific section of the artifact. You can also select pieces of text or code and ask AI to modify those. **Save your work:** Make sure that if you manually edit the artifact, you first save the artifact before you ask the AI to continue editing it, you can do this by clicking on the "save" icon on the top right of the artifact. When you ask the AI to edit the artifact a new version will automatically be saved. Previous versions will also be saved, so you can track changes and revisit previous iterations if necessary by clicking on the arrows at the top left of the artifact. The completed artifact can then be downloaded for use in your research paper, presentation, or other academic projects. Do this by clicking on the arrow at the top right corner of the artefact, and selecting "download as".

2.6 Study Mode

“Study Mode” in UvA AI Chat turns the assistant into a guiding tutor. Instead of giving answers right away, the AI first asks questions, offers hints, and waits for your response before continuing. It helps prevent you from outsourcing too much of your thinking to AI. By actively practicing reasoning and formulating your own answers, you keep training and developing your critical thinking skills. The goal is that you reason yourself and thus reach better understanding.

How to use Study Mode:

1. **Activate and set a learning goal:** Open the slide menu in the textbox of UvA AI Chat and switch on Study Mode. To turn it off, click the cross next to Study. Then set a learning goal by specifying what you want to learn about and mention what you already know. Examples:

“I am preparing a seminar on the French revolution and specifically the storming of the Bastille, help me understand the causes and interpretations of the cause.”

“I am preparing for an exam in research skills and want to understand how to compare two groups using statistics. Explain which methods are suitable, such as the t-test, and when it is better to use another test.”

2. **Work in dialogue:** The AI asks open questions and waits for your answer. You can indicate how you want to be guided: ask for smaller steps, hints instead of answers, or feedback on your reasoning. You can also request specific forms, such as a multiple-choice question, concrete examples, or metaphors.
3. **Adjust pace and level:** Redirect where needed: *“use simpler language”*, *“make it harder”*, or *“go faster/slower”*. The AI adjusts its questions, explanations, and examples accordingly.
4. **Conclude and record:** To consolidate your learning, finish with a recap or practice material. Examples: *“Summarise in five bullet points the most important lessons from this conversation”* or *“Give three practice questions with short solutions.”* If you want to save the result or work it out further, combine Study Mode with Artefacts. This allows you to edit notes, tables, or draft texts in a separate window and retrieve them later.

Part 3: Personalization and collaboration

3.1 Further configuration

To get the most out of UvA AI Chat, it's worth configuring it to match your preferences. To do this, go to the settings by clicking the person icon at the bottom left of the screen. Here, you can immediately set your preferred theme (light/dark). Next, click on 'Settings,' then on 'Personalization' for the most important personalization options:

- **'Memory Creation':** Enable this to let UvA AI Chat save information about your previous prompts and conversations. This allows the AI to remember context from earlier interactions, such as your field of study (e.g. astronomy) or your hobbies (like cooking).
- **'Memory Context':** For the chat to actually use your stored 'memories' in new conversations, you'll also need to enable this option.
- **'Memory Management':** Here, you can review and manage the information UvA AI Chat has stored about you. You can adjust or delete saved memories at any time.
- **'Custom Instructions':** Specify how you want the AI to behave or the writing style it should use by default. These instructions are applied in the background to every conversation you start (unless you're using a persona).

Practical example of 'Custom Instructions':

In the 'Custom Instructions' field, you might enter something like:

"Always answer in Dutch. Write your responses as an academic advisor: supportive, critical, and focused on improving my work. Use formal language and avoid excessive jargon. Structure complex answers with bullet points for clarity".

3.2 Customized interaction: working with personas

Personas are an advanced feature that let you instruct UvA AI Chat to take on a specific role or character during a conversation. This makes interactions more focused and the output better aligned with your context.

The basics: definition and purpose

Personas are predefined 'characters' or 'personalities' that the AI can adopt to provide directly relevant answers. By giving a persona clear instructions, you set its role, expertise, tone, background knowledge, and function—for example, as a critical reviewer, subject-matter expert, or helpful tutor. Personas are especially useful for recurring tasks or scenarios where you want to use the same prompts repeatedly. This ensures you get consistent, relevant answers without having to re-enter all your context and instructions each time. With personas, you can have the AI communicate from a specific perspective, expertise, or knowledge area. Essentially, a persona is a complex, pre-saved 'master prompt' that shapes the AI's behavior for an entire conversation.

Standard personas

UvA AI Chat comes with several standard personas you can use right away, including:

- **Course Manual Creator:** Collaborates with you to create a course manual.
- **Socratic Coach:** Uses the Socratic method to help you deepen your understanding of a topic without giving direct answers.
- **Technical Assistant:** Supports users with technical questions by providing clear explanations and practical solutions related to UvA AI Chat.
- **Language Assistant:** Checks texts and offers tips on grammar, spelling, and style to improve language skills and optimize usage.

Creating your own personas: a step-by-step guide

Building your own tailored personas can take your AI interactions to the next level. This lets you set up the persona's behavior exactly the way you want. Go to the mask icon in the left sidebar to access "Personas," then choose the option to create a new persona. Follow these steps:

1. **"Persona title":** Enter a short, recognizable name for your persona.
2. **"Persona description":** Briefly describe the role and purpose of the persona (this does not affect its behavior).
3. **"Which AI model do you want to use?":** Select the underlying language model (e.g., GPT-4o) for your persona. See section 4.1.
4. **"Choose a conversation style":** Pick a style. 'Balanced' is often a good starting point, but you can also select 'Creative' (open, imaginative responses) or 'Precise' (concise and factual).
5. **"Describe your persona and its behaviour":** This is the key step. Give a detailed description of the tone, expertise, and behavior you want your persona to have. You can adapt a preset or fill in this field manually—be as specific as possible.
6. **"Upload files to enhance your persona's resources":** Add specific documents, such as manuals or style guides, that your persona should use as a knowledge base.
7. **"Save":** Click "Save" to store and activate your persona.

Practical example of a custom persona:

A researcher wants an AI assistant that helps critically review her draft articles. She creates this persona:

- **Persona title:** "Critical Research Assistant".
- **Persona description:** "Analyzes academic texts for logical consistency, weak arguments, and unclear phrasing".
- **Describe your persona and its behaviour:** "You are an experienced peer reviewer for an academic journal. Your tone is critical but always constructive. When giving feedback on an uploaded text, structure your answer in three parts: 1) a brief summary of the author's main argument, 2) a list of strengths, and 3) a detailed list of points for improvement, including specific suggestions for rephrasing and questions that challenge the author to reflect on weaknesses in their reasoning".

3.3 “Groups”: collaborating and sharing with others

The “Groups” feature makes it easy to work together on shared projects. It’s ideal for teamwork—whether you’re conducting research, preparing a joint presentation, or working on any other project. Within a group, you can easily share files, see each other’s prompts, and work toward the same goals. Find the “Groups” function via the two-person icon in the left sidebar.

To create a group, click the “Add Group” button. Fill in the following textboxes:

Group Name: Enter a name for your group in the "Group Name" field. This is a required field.

Group Description: Provide a description for the group in the "Group Description" text box. This is generally the purpose of the group.

Members: Add the email addresses of the people you want to be members of the group. You can separate multiple email addresses with a comma.

Owners: Enter the email addresses of the people who will be the owners of the group. These should also be separated by commas. These owners will be able to edit the group.

Personas: If applicable, select the personas for the group.

Prompts: If applicable, select any specific saved prompts for the group.

Start and End Dates: Choose a start date and an end date for the group using the date pickers.

Save: Click the "Save" button to finalize the creation of the group.

By creating a group, you can share specific chats, personas, or prompts. You control whether added members can only use the personas (‘Members’) or also edit them (‘Owners’). As a teacher, for example, you might add students as members so they can use a specific persona you’ve created. You can also set a start and end date for the group if needed.

3.4 “My Projects”: your organized workspace

Under “My Projects” (the folder icon in the left sidebar), you can set up your own projects. This acts as a digital container for all materials related to a specific task or research project. Use it to keep your chats organized, especially if you have multiple conversations on the same topic. To get started, click ‘+ Add Project’ at the top right of the screen. Within a project, you can collect chats, prompts, and personas in one place. This lets you easily navigate back to earlier prompts and responses, making it simple to pick up right where you left off.

Part 4: Advanced usage

4.1 Choosing the right engine

UvA AI Chat gives you access to a range of advanced AI models (Large Language Models). The default, GPT-4o, is a strong all-rounder and works well for most tasks, but for specific needs, a different model may deliver better results. For simple tasks, you might prefer a smaller, more efficient model that uses less energy. The table below is a quick reference to help you choose the most suitable AI model for your task. To find out which model works best for your needs, you'll need to experiment - especially if you're running recurring, specific tasks.

Table 2: Comparison of available AI models

Model	General Use Cases	Knowledge Cutoff	Energy / Cost (relative)	Type	Context Window Input	Context Window Output
gpt-5	Explaining complex concepts, advanced coding projects, interpreting academic papers	30-09-2024	High	Advanced reasoning model	272K	128K
gpt-5-mini	moderate concept clarification, brainstorming ideas, study planning, understanding course material	31-05-2024	Medium	Efficient reasoning model	272K	128K
gpt-5-nano	Quick definitions, fact-checking, basic explanations, vocabulary help, concept reviews	31-05-2024	Low	Lightweight reasoning model	272K	128K
GPT-4o	Analyzing diagrams/charts, visual content interpretation, presentation feedback	20-11-2024	High	Multimodal model	128K	16K
GPT-4.1	General, language, high input, creative tasks, agentic planning	14-04-2025	High	Advanced language model	1M	100K
Llama-3.3-70B-Instruct AWQ	Open source, local deployment, long context	01-12-2023	Medium	Open source language model	128K	8K

4.2 Expanding functionality with extensions

Extensions are intended for technically savvy users who are familiar with APIs. They work as extra tools that the AI can use to handle tasks outside of the chat environment, such as retrieving information from external databases or carrying out actions in other software

How it works

Extensions are powerful tools that give UvA AI Chat greater capabilities by allowing the AI to make *API calls* to internal or external systems. They act as extra tools that enable the AI to perform tasks outside of the chat environment, such as retrieving information from a database, carrying out actions in other software (like adding an item to a to-do list), or sending and receiving data. These tools are intended for technically savvy users who are familiar with APIs, as incorrect use can trigger unintended actions in external systems.

The process involves defining the extension's details and functions, and it utilizes the API structure described in the official OpenAI documentation ([via openai.com](https://openai.com)). The creation interface is shown in the provided image.

To add your own extension, click “Add extension”

Name: Give your extension a name in the "Name of your Extension" field.

Short description: Write a brief description of the extension.

Detail description: Provide a more detailed explanation of its specialties and the steps required to execute it.

Headers: Define the necessary headers for the API calls. A default "Content-Type" header with the value "application/json" is shown. You can add more headers by clicking "Add Header." The platform also supports securing header values stored in Azure Key Vault.

Functions: Add the specific functions that the extension will perform by clicking "Add Function." These functions can support various API requests, including GET, POST, and PUT, allowing the extension to both retrieve data and trigger actions.

Submit: Once all the details are filled out, click the "Submit" button to finalize the creation of your extension.

Practical example of using an extension:

A researcher configures an extension that communicates with the UvA library catalog API. Now they can use a prompt like:

“Use the library extension to find the five most recent publications by author ‘Adriaan van Dis’. Provide the full APA citations for each publication and a direct link to each in the catalog.”

Frequently Asked Questions (FAQ)

What should I do if the AI gives strange or incorrect answers?

First, check your prompt: was it clear and specific enough? Try rephrasing your question. Sometimes it helps to start a new conversation or select a different AI model.

What's the difference between “Custom Instructions” and a “Persona”?

“Custom Instructions” are general directions you give the AI for all your conversations (unless a persona is active), like “Always answer in Dutch” or “Use a formal tone.” A “Persona” is much more specific—you give the AI a complete role (for example, “critical researcher” or “creative copywriter”) for a certain task or conversation. An active persona overrides your custom instructions.

What about privacy when I upload documents?

Any information you enter UvA AI Chat falls under the same security and privacy policies as your UvA emails and documents. It's not accessible to external parties and is not used for training AI. Of course, always use caution when uploading documents containing personal data of others, confidential business information, or other sensitive material. For more information, please consult our [policy framework and guidelines on GenAI in education](#).

The AI doesn't seem to fully understand my long conversations or documents. Why is that?

This probably relates to the “context window.” The AI can only “remember” a limited amount of text at a time. So, if you have a very long conversation or upload many documents, the AI may not be able to process all of it. If your conversation gets long, start a new chat (you can summarize the original chat and bring that summary with you). If your document is long, try asking specific questions about sections of it, summarize it yourself first, or break it up and let the AI summarize each part. You can also try selecting a model with a larger context window if one is available.

Why is the UvA developing its own AI infrastructure?

With UvA AI Chat, the University of Amsterdam offers a safe and responsible alternative to commercial generative AI (GenAI) tools. Developing its own AI infrastructure allows the UvA to:

- Maintain control over data (digital sovereignty);
- Safeguard privacy;
- Ensure equal access for staff and students;
- Support academic freedom;
- Provide transparency about the impact of the AI models used;
- Contribute to the conscious and responsible use of GenAI in education;
- Follow the leading core principles of the Dutch Code of Conduct on scientific integrity.

How is it different from ChatGPT, for example?

UvA AI Chat is similar to ChatGPT but is managed in-house and specifically developed for and by the UvA. This means all input remains under UvA control, making the university less dependent on big tech as an educational and research institution, and allowing the tool to be tailored to its specific needs.

A key feature is UvA's own front-end user interface, which provides control over how students, lecturers, and researchers interact with AI models. Through the implementation of an intermediary layer, UvA AI Chat is independent of any specific Large Language Model (LLM). The infrastructure supports various models, giving flexibility to choose the most appropriate one.

Who developed UvA AI Chat?

UvA AI Chat was developed by a multidisciplinary project team with colleagues from all faculties and services, as part of the project "Navigating the Ocean of Artificial Intelligence" (abbreviated as NOKI).

How does UvA AI Chat handle data and data protection?

All input is managed by the UvA. This means both input and output are protected, and the application can be used securely. Input from UvA AI Chat is not used to train or improve language models, so no model training takes place.

As an individual user, you can delete your chat history. Once deleted, your chat history cannot be restored.

Are there examples of successful UvA AI Chat implementations?

Yes, several pilot projects have been conducted using UvA AI Chat. You can learn more about these pilots through the UvA Teaching & Learning Centres. Check out a few of them [here](#).

Glossary

Prompt

Your instruction, question, or text for the AI. The foundation of everything

Token

The smallest unit of text for the AI (a word, or part of a word). The “context window” is measured in tokens.

Context Window

The ‘working memory’ of the AI—the maximum amount of text (tokens) the model can process at one time.

Persona

A set of instructions that gives the AI a specific role or style (e.g., expert, coach).

Hallucination

When the AI presents nonsense or wrong info that sounds convincing. Always double-check.

Prompt Engineering

The art of writing effective prompts to get exactly the output you want.

API (Application Programming Interface)

Like a power strip for software—lets different programs communicate. Extensions use this.

Extension

An extra tool or plug-in that gives UvA AI Chat new features, often by linking with other systems via an API.

Bias

When the AI (unintentionally) adopts biases from its training data, which can result in unfair or skewed answers.

LLM (Large Language Model)

A very large AI model trained on massive amounts of text, capable of understanding and generating human language. This is the engine behind the chat.